


# Process Synchronization





# *Critical Section Problem*

*At any given point in time, many kernel-mode processes may be active in OS, as a result kernel code is subject to race conditions.*

*Approaches to handle CS in OS:*

- ***Preemptive Kernel*** – *A preemptive kernel allows a process to be preempted while it is running in kernel mode.*
- ***Non-Preemptive Kernels*** – *A non-preemptive kernel doesn't allow a process running in kernel mode to be preempted.*





# *Synchronization Hardware*

---

- ❑ *Hardware features can make any programming task easier and improve efficiency.*
- ❑ **Interrupt prevention** *(no pre-emption) while a shared variable is being modified, could be a solution to critical section problem on single processor systems. Unfortunately this is not a solution on multiprocessor environment following problems could occur.*
  - Time consuming on multiprocessor.
  - System Efficiency will decrease.
  - System clock could be effected.





# *Synchronization Hardware*

---

**Locking:** *Protecting critical regions through the use of **locks**.*





# *Synchronization Hardware(test\_and\_set Instruction)*

---

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

- *Executed atomically.*
- *Returns original value of passed parameter.*
- *Set new value of passed parameter to TRUE.*

- *Boolean variable lock initialized to false*
- *If lock = false, return false, set lock =true*
- *If lock = true, return true, lock true*

```
do {
    while (test_and_set(&lock))
        ; /* do nothing */
        /* critical section */
    lock = false;
        /* remainder section */
} while (true);
```

## *Solution*



# *Synchronization Hardware (Compare\_and\_swap())*

```
int compare_and_swap(int *value, int expected, int
new_value) {
    int temp = *value;

    if (*value == expected)
        *value = new_value;
    return temp;
}
```

- *Boolean variable lock initialized to 0*

```
do {
    while (compare_and_swap(&lock, 0, 1) != 0)
        ; /* do nothing */
    /* critical section */
    lock = 0;
    /* remainder section */
} while (true);
```

*Solution*



## *Synchronization Hardware bounded waiting*

---

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = test_and_set(&lock);
    waiting[i] = false;
    /* critical section */
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = false;
    else
        waiting[j] = false;
    /* remainder section */
} while (true);
```





# *Synchronization Hardware bounded waiting*

## Entry section

```
waiting[i] = true;
key = true;
while (waiting[i] && key)
key = test and set(&lock);
waiting[i] = false;
```

## Critical Section



## Exit section

```
j = (i + 1) % n;
while ((j != i) && !waiting[j])
j = (j + 1) % n;
if (j == i)
lock = false;
else
waiting[j] = false;
```

Waiting [i]  
Key

P0	P1	P2	P3
<b>F</b>	<b>F</b>	<b>F</b>	<b>F</b>

lock



## TestAndSet

```
if lock=false, return false,lock=true
if lock=true,return true,lock=true
```







# *Hardware Solutions*

---

- *Complicated Solutions.*
- *Generally inaccessible to application programmers.*





# ***Mutex locks***

---

- ***Mutex lock:** It a software tool to solve critical section problem and it is taken from (**Mutual Exclusion**).*
- *Mutex locks are used to protect critical section and to prevent race conditions.*
- *A process must **acquire the lock** before entering a critical section, it **releases the lock** when it exits the critical section.*
- *The **acquire()** function will acquire the lock and **release()** function will release the lock.*





# *Mutex locks*

*do{*

acquire lock

*critical section*

release lock

remainder section

*}while(true);*

*acquire(){*

while (!*Available*);

*available =false;*

*}*

*release(){*

*available =true;*

*}*

*Calls to acquire and release should be atomic usually implemented by hardware instructions.*





## ***Mutex locks disadvantage***

---

- ***Busy waiting:** While a process is in its critical section any other process that tries to enter its critical section must loop continuously in the call to `acquire()`, this type of mutex lock is also called a spinlock because the process “spins” while waiting for lock to become available.*
- *This continual looping is clearly a problem in real multiprogramming systems where single CPU is shared among many processes.*
- *Busy waiting **waste CPU cycles** that some other process might be able to use productively.*





## ***Busy Waiting Advantage***

---

- ***Busy waiting/Spinlock:*** No context switch is required, when a process must wait on lock, and context may take considerable time.
- On multiprocessor systems one thread can spin on one processor while another thread performs its critical section on another processor.





# Semaphore

- Semaphore tools provides a more sophisticated way (than mutex locks) for process to **synchronize** their activities.
- Semaphore is a **data type**
  - Value (Integer)
  - Operations
- A **semaphore (S)** is an **integer variable** that, apart from initialization, is accessed only through two standard **atomic operations**: **wait(S)** and **signal(S)**.
- The wait() operation was originally termed as **P** and signal operation is termed as **V**.





# *Semaphore: Operations*

---

```
wait(S){  
    while (S<=0); // Busy Wait  
    S--;  
}
```

```
signal(S){  
    S++;  
}
```





# ***Semaphore: Applications:***

---

- *Semaphore are used to solve the critical section problem*
- *Process Synchronization*
- *Resource management.*







# *Types of Semaphore*

---

*There are two main types of semaphores:*

- ***Binary Semaphores:** This is also known as **mutex lock** and their value can range between 0 and 1.*
  - *It is used to implement the solution of critical section problem with multiple processes.*
- ***Counting Semaphores:** Its value can range over an unrestricted domain. It is used to control access to a resource that has multiple instances.*





# Types of Semaphore

---

- *Counting semaphores* are used to coordinate the resource access, where the semaphore count is the number of available resources.
- *Each process that wishes to use resource need to perform wait() on semaphore, when process release the resource it will perform signal() on semaphore.*
- *When count of the semaphore goes to 0 all resources are being used. After that wish to use a resource will block until count becomes greater than 0.*





# *Types of Semaphore*

---

- *Semaphore can be used to solve various synchronization problems.*

*Example: Consider two concurrently running processes: P1 with a statement S1 and P2 with a statement S2. Suppose we want that S2 should be execute only after S1 has completed.*

*S1:*  
*signal(S)*

*wait (S);*  
*S2*





# *Semaphores without busy waiting*

---

- *With each semaphore there is an associated waiting queue.*
- *Each entry in a waiting queue has two data items:*
  - *value (of type integer)*
  - *a pointer to the next record in the list*
- *Two operations:*
  - *block() – place the process invoking the operation on the appropriate waiting queue.*
  - *wakeup() – remove one of processes in the waiting queue and place it in the ready queue.*





# *Semaphores without busy waiting*

***Typedef struct{***

***int value;***

***Struct process \*list;*** }

***}semaphore;***

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}  
  
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```





# *Advantages of Semaphores*

---

*Some of the advantages of semaphores are as follows:*

- They follow the mutual exclusion principle strictly and are much more efficient than some other methods of synchronization.*
- There is no resource wastage because of busy waiting in semaphores as processor time is not wasted unnecessarily to check if a condition is fulfilled to allow a process to access the critical section.*





# Limitations of Semaphores

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes. Let  $S$  and  $Q$  be two semaphores initialized to 1

$P_0$	$P_1$
<code>wait(S);</code>	<code>wait(Q);</code>
<code>wait(Q);</code>	<code>wait(S);</code>
...	...
<code>signal(S);</code>	<code>signal(Q);</code>
<code>signal(Q);</code>	<code>signal(S);</code>

- **Starvation** – *indefinite blocking*
  - A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
  - Solved via **priority-inheritance protocol**





## *Semaphores Example - 1*

---

A shared variable  $x$ , initialized to zero, is operated on by four concurrent processes W, X, Y, Z as follows. Each of the processes W and X reads  $x$  from memory, increments by one, stores it to memory, and then terminates. Each of the processes Y and Z reads  $x$  from memory, decrements by two, stores it to memory, and then terminates. Each process before reading  $x$  invokes the P operation (i.e., wait) on a counting semaphore S and invokes the V operation (i.e., signal) on the semaphore S after storing  $x$  to memory. Semaphore S is initialized to two. What is the maximum possible value of  $x$  after all processes complete execution?

(A) -2

(B) -1

(C) 1

(D) 2 ←







Process P:

```
while (1) {  
W:  
    print '0';  
    print '0';  
X:  
}
```

Process Q:

```
while (1) {  
Y:  
    print '1';  
    print '1';  
Z:  
}
```

Synchronization statements can be inserted only at points W, X, Y and Z.

*Which of the following will always lead to an output starting with '001100110011'?*

- (A)  $P(S)$  at W,  $V(S)$  at X,  $P(T)$  at Y,  $V(T)$  at Z,  $S$  and  $T$  initially 1
- (B)  $P(S)$  at W,  $V(T)$  at X,  $P(T)$  at Y,  $V(S)$  at Z,  $S$  initially 1, and  $T$  initially 0
- (C)  $P(S)$  at W,  $V(T)$  at X,  $P(T)$  at Y,  $V(S)$  at Z,  $S$  and  $T$  initially 1
- (D)  $P(S)$  at W,  $V(S)$  at X,  $P(T)$  at Y,  $V(T)$  at Z,  $S$  initially 1, and  $T$  initially 0





# Bounded Buffer Problem

---

- *There is a buffer of  $n$  slots and each slot is capable of storing **one unit** of data. There are two processes running, namely, producer and consumer, which are operating on the buffer.*
- *A producer tries to **insert** data into an **empty** slot of the buffer. A consumer tries to remove data from a filled slot in the buffer. Two processes won't produce the expected output if they are being executed concurrently.*
- *There needs to be a way to make the producer and consumer work in an independent manner.*





# Solution of Bounded Buffer Problem

- **mutex**, a *binary semaphore* which is used to acquire and release the lock.
- **empty**, a *counting semaphore* whose initial value is the number of slots in the buffer, since, initially all slots are empty.
- **full**, a *counting semaphore* whose initial value is 0

The structure of the **producer** process

```
do {  
    // produce an item in nextp  
  
    wait (empty);  
    wait (mutex);  
  
    // add the item to the buffer  
  
    signal (mutex);  
    signal (full);  
  
} while (TRUE);
```

The structure of the **consumer** process

```
do {  
    wait (full);  
    wait (mutex);  
  
    // remove an item from  
    // buffer to nextc  
  
    signal (mutex);  
    signal (empty);  
  
    // consume the item in nextc  
} while (TRUE);
```



# Solution of Bounded Buffer Problem

---

The structure of the **producer** process

```
do {  
    // produce an item in nextp  
  
    wait (empty);  
    wait (mutex);  
  
    // add the item to the buffer  
  
    signal (mutex);  
    signal (full);  
  
} while (TRUE);
```

The structure of the **consumer** process

```
do {  
    wait (full);  
    wait (mutex);  
  
    // remove an item from  
    // buffer to nextc  
  
    signal (mutex);  
    signal (empty);  
  
    // consume the item in nextc  
} while (TRUE);
```



# Readers Writers Problem

---

*A data set is shared among a number of concurrent processes*

**Readers** – *only read the data set; they do not perform any updates*

**Writers** – *can both read and write.*

- **Problem** – *allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.*





# Solution of Readers Writers Problem

## *Shared Data:*

Semaphore *mutex* initialized to *1*.

Semaphore *wrt* initialized to *1*.

Integer *readcount* initialized to *0*.

The structure of a *writer* process

```
while (true) {  
    wait (wrt) ;  
    write operation  
    signal (wrt) ;  
}
```

The structure of a *reader* process

```
while (true) {  
    wait (mutex) ;  
    readcount ++ ;  
    if (readercount == 1)  
        wait (wrt) ;  
    signal (mutex)
```

## *Read Operation*

```
    wait (mutex) ;  
    readcount -- ;  
    if (redacount == 0)  
        signal (wrt) ;  
    signal (mutex) ;
```

```
}
```





# Dining-Philosophers Problem

- *There are five silent philosophers ( $P1 - P5$ ) sitting around a circular table, and they **eat and think alternatively**.*
- *There is a bowl of rice for each of the philosophers and 5 chopsticks ( $1 - 5$ ).*
- *To be able to eat, a **philosopher needs to have chopstick in both his hands**. A hungry philosopher may only eat if there are both chopsticks available.*
- *After eating a philosopher puts down their chopstick and begin thinking again.*





# Dining-Philosophers Problem Solution

---


- *A solution of the Dining Philosophers Problem is to use a semaphore to represent a chopstick.*
- *A chopstick can be picked up by executing a **wait** operation on the **semaphore** and released by executing a **signal** semaphore.*
- *The structure of the chopstick is shown below:*  
  
***semaphore chopstick [5];***
- *Initially the elements of the chopstick are initialized to 1 as the chopsticks are on the table and not picked up by a philosopher.*







# Dining-Philosophers Problem Solution



*semaphore chopstick [5];*

do {

*wait( chopstick[i] );*

*wait( chopstick[ (i+1) % 5] );*

.

*Eating the Rice*

*signal( chopstick[i] );*

*signal( chopstick[ (i+1) % 5] );*

*Thinking*

.

} while(1);



## Difficulty with the solution

*The above solution makes sure that no two neighboring philosophers can eat at the same time. But this solution can lead to a **deadlock**. This may happen if all the philosophers pick their left chopstick simultaneously. Then none of them can eat and deadlock occurs.*

*Some of the ways to **avoid deadlock** are as follows:*

- **There should be at most four philosophers on the table.***
- **An even philosopher should pick the right chopstick and then the left chopstick while an odd philosopher should pick the left chopstick and then the right chopstick.***
- **A philosopher should only be allowed to pick their chopstick if both are available at the same time.***

